



**Transformation**<sup>TM</sup>  
— LEADERS PODCAST —

## Qdrant Collections and Vectors: A Beginner's Guide

---

### What is Qdrant?

---

Qdrant is a vector database that stores and searches through "vectors" - lists of numbers that represent data like text, images, or any information you want to search through semantically.

Think of a Qdrant collection like a filing cabinet, but instead of storing documents, it stores vectors that represent the meaning of your data.

### Understanding Vectors

---

A vector is like a coordinate that describes something in multi-dimensional space.

**Simple Example:** Describing a person with numbers:

- Height: 5.8
- Weight: 150
- Age: 25

That's a 3-dimensional vector: [5.8, 150, 25]

**In Qdrant:** Vectors are usually much longer (100-1500+ numbers) and represent complex concepts like the "meaning" of a sentence or document.

### Creating Collections - The Basics

---

Creating a collection is straightforward:

```
client.create_collection(  
    collection_name="my_collection",  
    vectors_config=VectorParams(size=384, distance=Distance.COSINE)  
)
```

## Vector Parameters Explained

---

When creating a collection, you must specify three key parameters:

### 1. Size - How Many Numbers in Each Vector?

**Important:** You don't choose the size arbitrarily - your embedding model determines it.

Common embedding model sizes:

- OpenAI `text-embedding-3-small` : 1536 dimensions
- OpenAI `text-embedding-3-large` : 3072 dimensions
- Sentence Transformers `all-MiniLM-L6-v2` : 384 dimensions
- Google's models: Often 768 dimensions

**Rule:** The collection size must exactly match your embedding model's output size.

```
# If using OpenAI embeddings  
collection_config = VectorParams(  
    size=1536, # Must match OpenAI's output exactly  
    distance=Distance.COSINE  
)
```

**Smaller vs Larger Vectors:**

- **Smaller (384):** Faster searches, less storage, less detailed understanding
- **Larger (1536+):** More detailed understanding, slower searches, more storage

### 2. Distance - How to Measure Similarity

Three main options:

- **COSINE:** Best for text/semantic similarity (most common choice)
- **EUCLIDEAN:** Measures straight-line distance between points
- **DOT:** For when larger numbers indicate higher similarity

### 3. Vector Name (Optional)

You can store multiple types of vectors in one collection:

- "text\_vector" and "image\_vector" in the same collection
- Each can have different sizes and distance metrics

## Understanding PointStruct

---

`PointStruct` is Qdrant's way of organizing data - like a database row format:

```
PointStruct(  
  id=1,          # Unique identifier (like a primary key)  
  vector=[...],  # The embedding numbers from your model  
  payload={...}  # Your actual data and metadata  
)
```

## The Payload - Your Flexible Metadata

---

The payload can contain unlimited metadata about your data:

```
PointStruct(  
  id=1,  
  vector=embedding_vector, # From your embedding model  
  payload={  
    "review": "Great pizza!",  
    "rating": 5,  
    "customer_name": "John Smith",  
    "phone": "555-1234",  
    "visit_date": "2024-06-15",  
    "restaurant_name": "Tony's Pizza",  
    "location": "Downtown",  
    "price_range": "$$",  
    "cuisine_type": "Italian"  
  }  
)
```

## Complete Real-World Example

---

```

import openai
from qdrant_client import QdrantClient
from qdrant_client.models import PointStruct, VectorParams, Distance

# 1. Create collection (size matches OpenAI's model)
client.create_collection(
    collection_name="restaurant_reviews",
    vectors_config=VectorParams(size=1536, distance=Distance.COSINE)
)

# 2. Convert text to vector using OpenAI
review_text = "Great pizza, friendly service!"
embedding = openai.Embedding.create(
    input=review_text,
    model="text-embedding-3-small"
)
vector = embedding['data'][0]['embedding'] # This will be 1536 numbers

# 3. Store in Qdrant
client.upsert(
    collection_name="restaurant_reviews",
    points=[
        PointStruct(
            id=1,
            vector=vector, # The 1536 numbers from OpenAI
            payload={
                "review": review_text,
                "rating": 5,
                "customer": "Alice",
                "date": "2024-06-18",
                "restaurant": "Tony's Pizza",
                "location": "Downtown"
            }
        )
    ]
)

```

## The Power of Hybrid Search

---

Qdrant gives you two types of search capabilities:

### 1. Vector Search Only (Semantic Similarity)

Find content with similar meaning:

```
# Find reviews similar to "great food"
results = client.search(
    collection_name="restaurant_reviews",
    query_vector=embed_text("great food"),
    limit=5
)
```

## 2. Hybrid Search (Vector + Metadata Filtering)

Combine semantic search with precise filtering:

```
# Find reviews similar to "great food"
# BUT only from Italian restaurants in downtown with 4+ stars
results = client.search(
    collection_name="restaurant_reviews",
    query_vector=embed_text("great food"),
    query_filter=Filter(
        must=[
            FieldCondition(key="cuisine_type", match=MatchValue(value="Italian")),
            FieldCondition(key="location", match=MatchValue(value="Downtown")),
            FieldCondition(key="rating", range=Range(gte=4))
        ]
    ),
    limit=5
)
```

## Real-World Use Case Example

---

Building a restaurant recommendation system:

```
# User asks: "I want great pasta near downtown with good service"
search_query = "great pasta good service"
user_location = "downtown"

results = client.search(
    collection_name="restaurant_reviews",
    query_vector=embed_text(search_query), # Semantic similarity
    query_filter=Filter(
        must=[
            FieldCondition(key="location", match=MatchValue(value=user_location)),
            FieldCondition(key="rating", range=Range(gte=4))
        ]
    )
)

# Returns reviews that:
# 1. Are semantically similar to "great pasta good service" (vector search)
# 2. Are from downtown restaurants (metadata filter)
# 3. Have 4+ star ratings (metadata filter)
```

## Why Hybrid Search is Powerful

---

Traditional SQL databases can't understand that "delicious" and "tasty" mean similar things.

Pure vector search can't easily filter by exact criteria like "only Italian restaurants."

Qdrant's hybrid approach gives you:

- **Semantic understanding** from vectors ("great food" matches "delicious meal")
- **Precise filtering** from metadata (price range, location, date, ratings)
- **Fast searches** because Qdrant indexes both vectors and metadata

## Key Takeaways

---

1. **Vector size is determined by your embedding model** - you must match it exactly
2. **Only embed the meaningful text** (like review content) - store everything else as metadata
3. **Use COSINE distance for most text applications**
4. **Combine vector search with metadata filtering** for powerful, precise results
5. **Think of it as a semantic-aware SQL database** that understands meaning while allowing precise filtering

## Best Practices

---

- 
- Choose your embedding model first, then set collection size to match
  - Store rich metadata in the payload for flexible filtering
  - Use descriptive field names in your payload
  - Index frequently filtered fields for better performance
  - Test different distance metrics with your specific use case